

Main page

Science

Programming

Data and database

Web services

Google

Microsoft

Linux

Apple

Adobe

HTML5 & CSS

Python

C (programming language)

PHP, MySQL, Apache

JavaScript

Security

Software

Hardware

Gadget

Network Security Through Data Analysis: Building Situational Awareness (2014)

Part II. Tools

This section is about a number of tools for use in data analysis. The primary focus of this section is on two particular tools: SiLK and R. The System for Internet-Level Knowledge (SiLK) is a NetFlow analysis toolkit developed by the CERT at Carnegie Mellon University, which enables analysts to develop sophisticated flow analysis systems quickly and efficiently. R, a statistical analysis package developed at the University of Auckland, enables exploratory data analysis and visualization.

At this time, there is no killer app for network analysis. Analysis requires using many tools, often in ways they weren't really designed for. The tools covered in this section form what I believe to be a basic functional toolkit for an analyst. Combining them with a light scripting language such as Python empowers analysts to explore data and develop operationally useful products.

The remainder of this section is divided into five chapters. [Chapter 5](#) describes the SiLK suite, [Chapter 6](#) describes R. [Chapter 7](#) discusses IDS; while IDSes were briefly discussed in [Part I](#), this chapter discusses the construction and maintenance of these tools—analysts will often produce ad hoc IDSes to identify or deal with attacks. [Chapter 8](#) discusses tools to identify the ways in which hosts are connected to the Internet, including reverse DNS lookups, looking glasses, and tools such as *traceroute* and *ping*. Finally, [Chapter 9](#) discusses additional tools that are useful for particular analytic tasks.

Chapter 5. The SiLK Suite

SiLK, the System for Internet-Level Knowledge, is a toolkit originally developed by Carnegie Mellon's CERT to conduct large-scale netflow analysis. SiLK is now used extensively by the Department of Defense, academic institutions, and industry as a basic analytical toolkit.

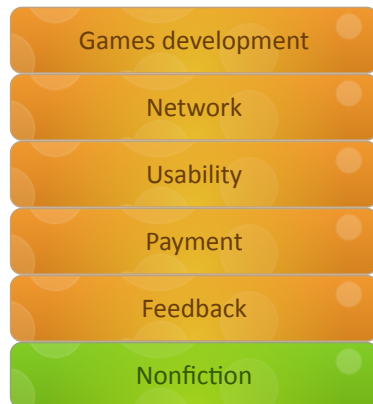
This chapter focuses primarily on using SiLK as an analytical tool. The CERT Network Situational Awareness team has published extensive [references](#) on using SiLK, installing collectors, and setting up the suite.

What Is SiLK and How Does It Work?

SiLK is a suite of tools for querying and analyzing NetFlow data. The SiLK suite enables an analyst to rapidly and efficiently query very large volumes of network traffic in order to identify complex aggregate phenomena or extract individual events.

Google Custom Search





SiLK is effectively a database at the command line. Each tool performs a specific query, manipulation, or aggregation of data, and commands are chained together to produce results. By chaining together multiple records along pipes, SiLK enables the analyst to create complex commands that field data along multiple channels simultaneously. For example, the following sequence of SiLK queries pull HTTP (port 80) traffic from flow data, producing a time series and a list of activity by busiest address. See [Example 5-1](#) for the basics of SiLK operation: commands are passed through a series of pipes, which can be stdin, stdout, or fifos (named pipes).

Example 5-1. Some overly complicated rfilter voodoo

```
$ mkfifo out2
```

```
$ rfilter --proto=6 --aport=80 data.rwf --pass=stdout |
```

```
    rfilter --input=stdin --proto=6 --pass=stdout
```

```
    --all=out2 | rwstats --top --count=10 --fields=1 &
```

```
    rwcoun out2 --bin-size=300
```

Data is maintained in an efficient binary representation up until the last moment, until commands that produce text (or some optional outputs) are called to produce output.

SiLK is very much an old-school Unix application suite: a family of tools tied together with pipes and using a lot of optional arguments. By using this approach, it's possible to create powerful analytic scripts with SiLK, because the tools have well-defined interfaces that will efficiently handle binary data. Effectively using SiLK involves connecting the appropriate tools together in order to process binary data and produce text only at the very end of the process.

This chapter also uses some basic Unix shell commands such as ls, cat, and head. I don't require you to know the shell on an expert level.

Acquiring and Installing SiLK

The SiLK homepage is maintained at the CERT NetSA Security Suite web page. The SiLK package is available free for download, and can be installed on most Unix systems without much difficulty. The CERT also provides a live CD image that can be used on its own.

The SiLK live CD comes with a training dataset called LBNL-05, anonymized header traces from Lawrence Berkeley National Labs in 2005. If you install the live CD, the data will be immediately accessible. If not, you can fetch the data from [The LBNL-05 reference data page](#).^[5]

In addition to the live CD, SiLK is available in several package managers, including homebrew.

The Datafiles

The LBNL datafiles are stored in a file hierarchy; [Example 5-2](#) shows the results of downloading and unarchiving them.

Example 5-2. Downloading the SiLK archives

```
$ gunzip -c SiLK-LBNL-05-noscan.tar
$ gunzip -c SiLK-LBNL-05-scanners.tar
$ cd SiLK-LBNL-05
$ ls
README-S0.txt  in      out      silk.conf
README-S1.txt  inweb   outweb
$ ls in/2005/01/07/*.01
in/2005/01/07/in-S0_20050107.01 in/2005/01/07/in-S1_20050107.01
```

When collecting data, SiLK partitions the data into subdirectories that divide traffic by the type of traffic and the time the event occurred. This provides scalability and speeds up analysis. However, it's also generally a black box, and one we're breaking right now simply to have some files to work with. For the purposes of demonstration and education, we're going to work with four specific files:

```
$ inweb/2005/01/06/iw-S0_20050106.20
$ inweb/2005/01/06/iw-S0_20050106.21
$ in/2005/01/07/in-S0_20050107.01
$ in/2005/01/07/in-S1_20050107.01
```

These files are not special in any way. I chose them just to provide examples of scan and nonscan traffic. The following data discusses how to partition data and what the filenames mean.

Choosing and Formatting Output Field Manipulation: `rwcut`

SiLK records are stored in a compact binary format. They can't be read directly, and are instead accessed using the `rwcut` tool (see [Example 5-3](#)). In the following example, and any other examples with an output greater than 80 characters, the lines are manually broken for clarity.

Example 5-3. Simple file access with `rwcut`

```
$ rwcut inweb/2005/01/06/iw-S0_20050106.20 | more
      sIP|      dIP|sPort|dPort|pro| packets| bytes|\
      flags|      sTime| dur|      eTime|sen|
148.19.251.179| 128.3.148.48| 2497| 80| 6| 16| 2631|\
FS PA [2005/01/06T20:01:54.119| 0.246|2005/01/06T20:01:54.365| ?|
148.19.251.179| 128.3.148.48| 2498| 80| 6| 14| 2159|\
S PA [2005/01/06T20:01:54.160| 0.260|2005/01/06T20:01:54.420| ?|
...
```

In its default invocation, `rwcut` outputs 12 fields: source and destination IP addresses and ports, protocol, number of packets, number of bytes, TCP flags, start time, duration, end time, and sensor of a flow. These values have been discussed previously in [Chapter 2](#), except for the sensor field. SiLK can be configured to identify individual sensors, which is useful when you're trying to figure out where traffic came from or where it's going. The sensor field is whatever ID is assigned during configuration. In the default data there are no sensors, so the value is set to a question mark (?).

All SiLK commands have built-in documentation. Typing `rwcut --help` brings up an enormous help page. We will cover the basic options. A fuller description of options can be found in the [SiLK documentation for `rwcut`](#).

The most commonly used `rwcut` commands select the fields displayed during invocation. `rwcut` can actually print 29 different fields, in arbitrary order. A list of these fields is in [Table 5-1](#).

`rwcut` fields are specified using the `--fields=` option, which takes the numeric values in [Table 5-1](#) or the string values, and prints the requested fields in the order specified, as in [Example 5-4](#).

Table 5-1. `rwcut` fields

Field	Numeric ID	Description
sIP	1	Source IP address
dIP	2	Destination IP address
sPort	3	Source port
dPort	4	Destination Port: if ICMP, the ICMP type and code is encoded here also
protocol	5	Layer 3 protocol
packets	6	Packets in the flow
bytes	7	Bytes in the flow
flags	8	OR of TCP flags
sTime	9	Start time in seconds
eTime	10	End time in seconds
dur	11	Duration (eTime-sTime)

Field	Numeric ID	Description
sensor	12	Sensor ID
in	13	SNMP ID of the incoming interface on the router
out	14	SNMP ID of the outgoing interface on the router
nhIP	15	Next hop address
sType	16	Classification of the source address (internal, external)
dType	17	Classification of the destination address (internal, external)
scc	18	Country code of the source IP
dcc	19	Country code of the destination IP
class	20	Class of the flow
type	21	Type of the flow
sTime +msec	22	sTime in milliseconds
eTime +msec	23	eTime in milliseconds
dur +msec	24	duration msec
icmpTypeCode	25	ICMP type and code
initialFlags	26	Flags in the first TCP packet
sessionFlags	27	Flags in all packets <i>except</i> the first
attributes	28	Attributes of the flow observed by the generator

Field	Numeric ID	Description
application	29	Guess as to the application in the flow

Example 5-4. Some examples of field ordering

\$# Show a limited set of fields

```
$ rwcut --field=1-5 inweb/2005/01/06/iw-S0_20050106.20 | head -2
```

```
    sIP|      dIP|sPort|dPort|pro|
```

```
148.19.251.179| 128.3.148.48| 2497| 80| 6|
```

\$#Note the -, now explicitly enumerate

```
$ rwcut --field=1,2,3,4,5 inweb/2005/01/06/iw-S0_20050106.20 | head -2
```

```
    sIP|      dIP|sPort|dPort|pro|
```

```
148.19.251.179| 128.3.148.48| 2497| 80| 6|
```

\$#Field order is based on what you enter in --field

```
$ rwcut --field=5,1,2,3,4 inweb/2005/01/06/iw-S0_20050106.20 | head -2
```

```
pro|      sIP|      dIP|sPort|dPort|
```

```
6| 148.19.251.179| 128.3.148.48| 2497| 80|
```

\$#We can use text instead of numbers

```
$ rwcut --field=sIP,dIP,proto inweb/2005/01/06/iw-S0_20050106.20 |head -2
```

```
    sIP|      dIP|proto|
```

```
148.19.251.179| 128.3.148.48| 6|
```

rwcut supports a number of other output formatting and manipulation tools. Some particularly useful ones, which let you control the lines that appear in the output, include:

--no-title

Commonly used with SiLK commands that produce tabular output. Drops the title from the output table.

--num-recs

Outputs a specific number of records, eliminating the need for the head pipe in the previous example. The default value is zero, which makes rwcut dump the entire contents of whatever file it's reading.

--start-rec-num *and* --end-rec-num

Can be used to fetch a range of records in the file.

Example 5-5 shows a few ways to manipulate record numbers and headers.

Example 5-5. Manipulating record numbers and headers

\$# Drop the title

```
$ rwcut --field=1-9 --no-title inweb/2005/01/06/iw-S0_20050106.20 | head -5
```

```
148.19.251.179| 128.3.148.48| 2497| 80| 6|    16|    2631|FS PA
```

```
|2005/01/06T20:01:54.119|
```

```
148.19.251.179| 128.3.148.48| 2498| 80| 6|    14|    2159| S PA
```

```
|2005/01/06T20:01:54.160|
```

```
148.19.251.179| 128.3.148.48| 2498| 80| 6|     2|     80|F  A
```

```
|2005/01/06T20:07:07.845|
```

```
56.71.233.157| 128.3.148.48|48906| 80| 6|     5|    300| S
```

```
|2005/01/06T20:01:50.011|
```

```
56.96.13.225| 128.3.148.48|50722| 80| 6|     6|    360| S
```

```
|2005/01/06T20:02:57.132|
```

\$# Drop the head statement

```
$ rwcut --field=1-9 inweb/2005/01/06/iw-S0_20050106.20 --num-recs=5
```

```
      sIP|      dIP|sPort|dPort|pro|  packets|  bytes|  flags
```

```
|      sTime|
```

```
148.19.251.179| 128.3.148.48| 2497| 80| 6|    16|    2631|FS PA
```

```
|2005/01/06T20:01:54.119|
```

```
148.19.251.179| 128.3.148.48| 2498| 80| 6|    14|    2159| S PA
```

```
|2005/01/06T20:01:54.160|
```

```
148.19.251.179| 128.3.148.48| 2498| 80| 6|     2|     80|F  A
```

```
|2005/01/06T20:07:07.845|
```

```
56.71.233.157| 128.3.148.48|48906| 80| 6|     5|    300| S
```

```
|2005/01/06T20:01:50.011|
```

```
56.96.13.225| 128.3.148.48|50722| 80| 6|     6|    360| S
```

```
|2005/01/06T20:02:57.132|
```

```
$# Print only the third through fifth record
```

```
$ rwcut -field=1-9 inweb/2005/01/06/iw-S0_20050106.20 --start-rec-num=3
```

```
--end-rec-num=5
```

```
      sIP|      dIP|sPort|dPort|pro|  packets|  bytes|  flags
|      sTime|
```

```
148.19.251.179| 128.3.148.48|2498| 80| 6|    2|    80|F  A
```

```
|2005/01/06T20:07:07.845|
```

```
56.71.233.157| 128.3.148.48|48906| 80| 6|    5|   300| S
```

```
|2005/01/06T20:01:50.011|
```

```
56.96.13.225| 128.3.148.48|50722| 80| 6|    6|   360| S
```

```
|2005/01/06T20:02:57.132|
```

A number of options manipulate output format. Tabulation is controllable with the `--column-separator`, `--no-final-column`, and `--no-columns` switches. `--column-separator` will change the character used to distinguish columns, while `--no-final-column` drops the delimiter at the end of the line. `--no-columns` removes any space padding between columns. The `--delimited` switch combines all three: it takes a character as an argument, uses that character as a column separator, removes all padding in the columns, and drops the final column separator.

In addition, there are a variety of switches for changing column content:

```
--integer-ips
```

Converts IP addresses to integers rather than dotted quads. This switch is deprecated as of SiLK v3, and users should now use `--ip-format=decimal`.

```
--ip-format
```

The updated version of `--integer-ips`, `--ip-format` specifies how addresses are rendered. Options include canonical (dotted quad for IPv4, canonical IPv6 for IPv6), zero-padded (canonical, except zeroes are expanded to the maximal value for each format, so 127.0.0.1 is 127.000.000.001), decimal (print as the corresponding 32-bit or 128-bit integer), hexadecimal (print the integer in hexadecimal format), and `force-ipv6` (prints all addresses in canonical IPv6 format, including IPv4 addresses mapped to the `::ffff:0:0/96` netblock).

```
--epoch-time
```

Prints timestamps as epoch values with floating-point millisecond precision.

```
--integer-tcp-flags
```

Converts TCP flags to their integer equivalents.

```
--zero-pad-ips
```


Pads the dotted quad IP address format with zeros, so that 128.2.11.12 is printed as 128.002.011.012. Depreciated in favor of `--ip-format` in SiLK v3.

`--icmp-type-and-code`

Places the ICMP type in the source port and the ICMP code in the destination port.

`--pager`

Specifies the program to use for paging output.

Example 5-6 shows some of the preceding options.

Example 5-6. Other formatting examples

\$# Change from fixed with columns to delims

```
$ rwcut --field=1-5 inweb/2005/01/06/iw-S0_20050106.20 --no-columns --num-recs=2
```

```
sIP|dIP|sPort|dPort|protocol|
```

```
148.19.251.179|128.3.148.48|2497|80|6|
```

```
148.19.251.179|128.3.148.48|2498|80|6|
```

\$# Change the column separator

```
$ rwcut --field=1-5 inweb/2005/01/06/iw-S0_20050106.20 --column-sep=:
```

```
--num-recs=2
```

```
sIP:      dIP:sPort:dPort:pro:
```

```
148.19.251.179: 128.3.148.48: 2497: 80: 6:
```

```
148.19.251.179: 128.3.148.48: 2498: 80: 6:
```

\$# Use `--delim` to change everything at once

```
$ rwcut --field=1-5 inweb/2005/01/06/iw-S0_20050106.20 --delim=: --num-recs=2
```

```
sIP:dIP:sPort:dPort:protocol
```

```
148.19.251.179:128.3.148.48:2497:80:6
```

```
148.19.251.179:128.3.148.48:2498:80:6
```

\$# Convert IP addresses to integers

```
$ rwcut --field=1-5 inweb/2005/01/06/iw-S0_20050106.20 --integer-ip --num-recs=2
```

```
sIP|    dIP|sPort|dPort|pro|
```

```
2484337587|2147718192| 2497| 80| 6|
```

```
2484337587|2147718192| 2498| 80| 6|
```

```
$# Use epoch time
```

```
$ rwcut --field=1-5,9 inweb/2005/01/06/iw-S0_20050106.20 --epoch --num-recs=2
```

```
      sIP|      dIP|sPort|dPort|pro|      sTime|
```

```
148.19.251.179| 128.3.148.48| 2497| 80| 6|1105041714.119|
```

```
148.19.251.179| 128.3.148.48| 2498| 80| 6|1105041714.160|
```

```
$# Zero pad IP addresses
```

```
$ rwcut --field=1-5,9 inweb/2005/01/06/iw-S0_20050106.20 --zero-pad --num-recs=2
```

```
      sIP|      dIP|sPort|dPort|pro|      sTime|
```

```
148.019.251.179|128.003.148.048| 2497| 80| 6|2005/01/06T20:01:54.119|
```

```
148.019.251.179|128.003.148.048| 2498| 80| 6|2005/01/06T20:01:54.160|
```

You will note that, as the command lines get more complex, I have truncated the longer options. SiLK uses GN U-style long options universally, so the only requirement for specifying an option is to type enough characters to make the name unambiguous. Expect more and more truncation as we build more and more complex commands.

Basic Field Manipulation: **rwfilter**

The most basic SiLK command with analytical values is `rwcut` paired with `rwfilter` through a pipe. [Example 5-7](#) shows a simple `rwfilter` command.

Example 5-7. A simple `rwfilter` command

```
$ rwfilter --dport=80 inweb/2005/01/06/iw-S0_20050106.20 --pass=stdout
```

```
| rwcut --field=1-9 --num-recs=5
```

```
      sIP|      dIP|sPort|dPort|pro| packets| bytes| flags
```

```
|      sTime|
```

```
148.19.251.179| 128.3.148.48| 2497| 80| 6|    16|   2631|FS PA
```

```
|2005/01/06T20:01:54.119|
```

```
148.19.251.179| 128.3.148.48| 2498| 80| 6|    14|   2159| S PA
```

```
|2005/01/06T20:01:54.160|
```

```
148.19.251.179| 128.3.148.48| 2498| 80| 6|     2|    80|F  A
```

```
|2005/01/06T20:07:07.845|
```

```
56.71.233.157| 128.3.148.48|48906| 80| 6|     5|   300| S
```

```
|2005/01/06T20:01:50.011|
```

```
56.96.13.225| 128.3.148.48|50722| 80| 6| 6| 360| S
```

```
|2005/01/06T20:02:57.132|
```

rwfilter with a single filter (the --dport option in this case), and a single redirect (the --pass=stdout) is about as simple as you can get. rwfilter is the workhorse of the SiLK suite: it reads input (directly from a file, using a set of globbing specifications, or through a pipe), applies one or more filters to each record in the data, and then redirects the records based on whether a record matches the filters (passes) or doesn't match (fails).

SiLK's rwfilter [documentation](#) is humongous, but primarily consists of repetitively describing the filter specifications for every field, so don't be intimidated. rwfilter options basically do one of three things: they specify how to filter data, how to read data, or how to direct the result of those filters.

Ports and Protocols

The easiest filters to start with are --sport, --dport, and --protocol. As the names imply, they filter on the source port, destination port, and protocol, respectively (see [Example 5-8](#)). These values can filter on a specific value (e.g., --sport=80 will pass any traffic where the source port is 80), or a range specified with a dash or commas (so --sport=79-83 will pass anything where the source port is between 79 and 83 inclusive, and could be expressed as --sport=79,80,81,82,83).

Example 5-8. Example filtering on sport

```
$ rwfilter --dport=4350-4360 inweb/2005/01/06/iw-S0_20050106.20
```

```
--pass=stdout | rwcut --field=1-9 --num-recs=5
```

```
      sIP|      dIP|sPort|dPort|pro| packets| bytes| flags
|      sTime|
218.131.115.42| 131.243.105.35| 80| 4360| 6| 2| 80|F A
|2005/01/06T20:24:21.879|
148.19.96.160|131.243.107.239| 80| 4350| 6| 27| 35445|FS PA
|2005/01/06T20:59:42.451|
148.19.96.160|131.243.107.239| 80| 4352| 6| 4| 709|FS PA
|2005/01/06T20:59:42.507|
148.19.96.160|131.243.107.239| 80| 4351| 6| 15| 16938|FS PA
|2005/01/06T20:59:42.501|
148.19.96.160|131.243.107.239| 80| 4353| 6| 4| 704|FS PA
|2005/01/06T20:59:42.544|
```

```
$ rwfilter --sport=4000- inweb/2005/01/06/iw-S0_20050106.20
```

```
--pass=stdout | rwcut --field=1-9 --num-recs=5
```

```

      sIP|      dIP|sPort|dPort|pro| packets| bytes| flags
|      sTime|
56.71.233.157| 128.3.148.48|48906| 80| 6|    5|   300| S
|2005/01/06T20:01:50.011|
56.96.13.225| 128.3.148.48|50722| 80| 6|    6|   360| S
|2005/01/06T20:02:57.132|
56.96.13.225| 128.3.148.48|50726| 80| 6|    6|   360| S
|2005/01/06T20:02:57.432|
58.236.56.129| 128.3.148.48|32621| 80| 6|    3|   144| S
|2005/01/06T20:12:10.747|
56.96.13.225| 128.3.148.48|54497| 443| 6|    6|   360| S
|2005/01/06T20:09:30.124|

```

```
$ rwfilter --dport=4350,4352 inweb/2005/01/06/iw-S0_20050106.20
```

```
--pass=stdout | rwcut --field=1-9 --num-recs=5
```

```

      sIP|      dIP|sPort|dPort|pro| packets| bytes| flags
|      sTime|
148.19.96.160|131.243.107.239| 80| 4350| 6|   27|  35445|FS PA
|2005/01/06T20:59:42.451|
148.19.96.160|131.243.107.239| 80| 4352| 6|    4|   709|FS PA
|2005/01/06T20:59:42.507|
148.19.96.160|131.243.107.239| 80| 4352| 6|    1|    40| A
|2005/01/06T20:59:42.516|

```

```
$ rwfilter --proto=1 in/2005/01/07/in-S0_20050107.01 --pass=stdout
```

```
| rwcut --field=1-6 --num-recs=2
```

```

      sIP|      dIP|sPort|dPort|pro| packets|
35.223.112.236| 128.3.23.93| 0| 2048| 1|    1|
62.198.182.170| 128.3.23.81| 0| 2048| 1|    1|

```

```
$ rfilter --proto=1,6,17 in/2005/01/07/in-S0_20050107.01 --pass=stdout
```

```
| rwcut --num-recs=2 --fields=1-6
```

```
    sIP|      dIP|sPort|dPort|pro|  packets|
```

```
116.66.41.147|131.243.163.201| 4283| 1026| 17|    1|
```

```
116.66.41.147|131.243.163.201| 3131| 1027| 17|    1|
```

```
$ rfilter --proto=1,6,17 in/2005/01/07/in-S0_20050107.01 --fail=stdout
```

```
| rwcut --num-recs=2 --fields=1-6
```

```
    sIP|      dIP|sPort|dPort|pro|  packets|
```

```
57.120.186.177| 128.3.26.171|  0|  0| 50|    70|
```

```
57.120.186.177| 128.3.26.171|  0|  0| 50|    81|
```

Note the use of `--fail` in the last example. Because there are 255 potential protocols, specifying “everything but TCP, ICMP, and UDP” could be expressed in two ways: either by specifying everything you want (`--proto=0,2-5,7-16,18-`), or by using the `--fail` option. I’ll discuss more advanced manipulation of `--pass` and `--fail` in the next chapter.

Size

Volume (size) options (bytes and packets) are similar to the protocol and port options in that you express them numerically. Unlike the enumerations (ports and protocol), these numeric values can be expressed only as single digits or ranges, *not* as comma-separated values. So, `--packets=70-81` is acceptable, but `--bytes=1,2,3,4` is not.

IP Addresses

The simplest form of IP address filtering simply expresses the IP address directly (see [Example 5-9](#)). The following examples show strict filtering on the source (`--saddress`) and destination (`--daddress`) address, and the `--any-address` option. `--any-address` will match *either* source or destination addresses.

Example 5-9. Filtering on IP addresses

```
$ rfilter --saddress=197.142.156.83 --pass=stdout
```

```
in/2005/01/07/in-S0_20050107.01 | rwcut --num-recs=2
```

```
    sIP|      dIP|sPort|dPort|pro|  packets|  bytes|  flags|
```

```
    sTime|  dur|      eTime|sen|
```

```
197.142.156.83| 224.2.127.254|44510| 9875| 17|    12|  7163|    |
```

```
2005/01/07T01:24:44.359| 16.756|2005/01/07T01:25:01.115| ?|
```

```
197.142.156.83| 224.2.127.254|44512| 9875| 17|    4|  2590|    |
```

```
2005/01/07T01:25:02.375| 5.742|2005/01/07T01:25:08.117| ?|
```

```
$ rfilter --daddress=128.3.26.249 --pass=stdout
```

```
in/2005/01/07/in-S0_20050107.01 | rwcut --num-recs=2
```

```

sIP|      dIP|sPort|dPort|pro| packets|  bytes|  flags|
sTime|   dur|      eTime|sen|
211.210.215.142| 128.3.26.249| 4068| 25| 6|    7|    388|FS PA |
2005/01/07T01:27:06.789| 5.052|2005/01/07T01:27:11.841| ?|
203.126.20.182| 128.3.26.249|51981| 4587| 6|    56|   2240|F A |
2005/01/07T01:27:04.812| 18.530|2005/01/07T01:27:23.342| ?|

```

```
$ rwfilter --any-address=128.3.26.249
```

```
--pass=stdout in/2005/01/07/in-S0_20050107.01 | rwcut --num-recs=2
```

```

sIP|      dIP|sPort|dPort|pro| packets|  bytes|  flags|
sTime|   dur|      eTime|sen|
211.210.215.142| 128.3.26.249| 4068| 25| 6|    7|    388|FS PA |
2005/01/07T01:27:06.789| 5.052|2005/01/07T01:27:11.841| ?|
203.126.20.182| 128.3.26.249|51981| 4587| 6|    56|   2240|F A |
2005/01/07T01:27:04.812| 18.530|2005/01/07T01:27:23.342| ?|

```

Address options accept a variety of range descriptors. Each quad in an IP address can be expressed using the same comma-dash format that protocols and ports use. IP addresses will also accept the character *x* to mean 0-255. This expression can be used within each quad; SiLK will match each quad separately. In addition to this comma-dash format, SiLK can match on CIDR blocks.

SiLK supports IPv6 by using IPv6's colon-based notation. The following are all examples of valid IPv6 filters in SiLK, and [Example 5-10](#) shows how to filter them:

```
::ffff:x
```

```
::ffff:0:aaaa,0-5
```

```
::ffff:0.0.5-130,1,255.x
```

Example 5-10. Filtering IP ranges

```
$#Filtering on the last quad
```

```
$ rwfilter --daddress=131.243.104.x inweb/2005/01/06/iw-S0_20050106.20
```

```
--pass=stdout | rwcut --field=1-5 --num-recs=5
```

```

sIP|      dIP|sPort|dPort|pro|
150.52.105.212|131.243.104.181| 80| 1262| 6|

```

```
150.52.105.212|131.243.104.181| 80| 1263| 6|
```

```
59.100.39.174| 131.243.104.27| 80| 3188| 6|
```

```
59.100.39.174| 131.243.104.27| 80| 3191| 6|
```

```
59.100.39.174| 131.243.104.27| 80| 3193| 6|
```

Filtering a range of specific values in the third quad

```
$ rfilter --daddress=131.243.104,107,219.x inweb/2005/01/06/iw-S0_20050106.20
```

```
--pass=stdout | rwcut --field=1-5 --num-recs=5
```

```
sIP|      dIP|sPort|dPort|pro|
```

```
208.122.23.36|131.243.219.201| 80| 2473| 6|
```

```
205.233.167.250|131.243.219.201| 80| 2471| 6|
```

```
58.68.205.40| 131.243.219.37| 80| 3433| 6|
```

```
208.233.181.122| 131.243.219.37| 80| 3434| 6|
```

```
58.68.205.40| 131.243.219.37| 80| 3435| 6|
```

Using CIDR blocks

```
$ rfilter --saddress=56.81.0.0/16 inweb/2005/01/06/iw-S0_20050106.20
```

```
--pass=stdout | rwcut --field=1-5 --num-recs=5
```

```
sIP|      dIP|sPort|dPort|pro|
```

```
56.81.19.218|131.243.219.201| 80| 2480| 6|
```

```
56.81.16.73|131.243.219.201| 80| 2484| 6|
```

```
56.81.16.73|131.243.219.201| 80| 2486| 6|
```

```
56.81.30.48|131.243.219.201| 443| 2490| 6|
```

```
56.81.31.159|131.243.219.201| 443| 2489| 6|
```

Time

There are three time options: `--stime`, `--etime`, and `--active-time`. These fields require a time range, which in SiLK is written in the format:

```
YYYY/MM/DDTHH:MM:SS-YYYY/MM/DDTHH:MM:SS
```

Note the T separating the day and hour. The `--stime` and `--etime` fields filter exactly what it says on the can, which can be a bit counterintuitive; specifying `--stime=2012/11/08T00:00:00-2012/11/08T00:02:00` filters any record whose *start time* is between midnight and two minutes after midnight on November 8, 2012. Records that

started *before* midnight and are still being transmitted during that range will not pass. To find records that occurred within a particular period, use the `--active-time` filter.

TCP Options

Flows are aggregates of packets, and in the majority of cases, this aggregation is relatively easy to understand. For example, the number of bytes in a flow is the sum of the number of bytes of all the packets that comprise the flow. TCP flags, however, are a bit more problematic. In NetFlow v5, a flow's flags are the bitwise OR of the flags in its constituent packets—meaning that a flow indicates that a flag was present or absent in the *entire* flow, but not *where*. A flow could conceivably consist of a gibberish sequence of flags such as a FIN, then an ACK and SYN. Monitoring software such as YAF expands NetFlow to include additional flag fields, which SiLK can take advantage of.

The core flag filtering switches are `--flags-initial`, `--flags-all`, and `--flags-session`. These options accept flags in the form *high flags/mask flags*. If a flag is listed in the mask, SiLK always parses it. If a flag is listed in the high flags, SiLK passes it *only* if the value is high. The flags themselves are expressed using the characters in [Table 5-2](#).

Table 5-2. Expressing TCP flags in *rwfilter*

Character	Flag
F	FIN
S	SYN
R	RST
P	PSH
A	ACK
U	URG
E	ECE
C	CWR

The combination of high flags and mask flags tends to confuse people, so let's review some examples. Remember that the basic rule is that in order to evaluate a flag, it *must* be in the mask. A flag specified as high but not specified in the mask will be ignored.

§ Setting the value to S/S will pass any record where the SYN flag is high.

§ Setting the value to S/SA will pass any record where the SYN flag is high *and* the ACK flag is low.

§ Setting the value to SA/SA will pass any record where *both* SYN and ACK flags are high.

§ A combination like SAF/SAFR will return any record where the SYN, ACK, FIN flags are high *and* the RST flag is low, which would be expected of a normal TCP connection.

In addition to these options, SiLK provides a set of flag-specific options in the form of --syn-flag, --fin-flag, and so on for each potential flag. These options take a 1 or 0 as an argument: setting the value to 1 will pass records where the flag is high, 0 will pass records where the flag is low, and not including the option will pass all records.

WHAT SHOULD TCP FLAGS LOOK LIKE?

The combination of TCP flags in any particular flow can be a useful indicator of the flow's behavior, and there are certain flag combinations that raise suspicion.

Almost all TCP flows should pass *either* SAF/SAFR or SAR/SAFR, *without* passing SAFR/SAFR. This is because most sessions will end in a FIN, with aberrations ending in a RST. If both FIN and RST are seen, that's suspicious.

A TCP session without an ACK flag is curious, *especially* if that session has four or more packets. Stacks are usually hardcoded to give up after *n* packets, where *n* tends to be in the neighborhood of three.

For a client, the initial flag should be a SYN, while a server should have a SYN+ACK. You should never see a SYN after the initial flag. Resynchronization would mean a new session started using the same ephemeral port, which is weird for TCP.

The PSH and URG flags are, in my mind, the universal indicator of boring sessions. If I see a session *without* PSH, especially if the session is long, it strikes me as curious. In my mind, a "normal" TCP session will have FSPA high. A flow with just PA high is usually a keep-alive and an indication of a broken flow—look in the repository for the same address combination and you'll probably find a SAP flow occurring before it.

Backscatter/response messages include A, SA, and RA flows. A good number of RA packets will arrive on any large network due to backscatter from spoofed DDoS attacks. There isn't really anything you can do about these packets; they're not even directly aimed at your network.

Helper Options

If you compare rfilter's option-based filtering against tcpdump's BPF filtering, it's immediately obvious that rfilter's approach is much more primitive. This was an intentional decision: rfilter is focused on processing large volumes as quickly as possible, and the overhead involved in processing some kind of parseable language was deemed too expensive.

The place where this usually trips people up is the lack of obvious not and or operators. For example, if you want to filter out all web sessions, you may try to filter traffic where one port is 80, and the other is ephemeral. The initial attempt might be:

```
rfilter --sport=80,1024-65535 --dport=80,1024-65535 --pass=stdout
```

The problem is that this will also pass any flows where the source and destination port are both 80, and flows where the source and destination port are both ephemeral. To deal with problems like this, rfilter has a collection of helper functions, which combined with the --fail option and multiple filters should be able to address any of these problems.

In the case of ports, the --aport option refers to either the source *or* the destination port. Using --aport and two filters, you can identify the appropriate sessions as follows:

```
rwfilter --aport=80 --pass=stdout | rwfilter --input-pipe=stdin
```

```
--aport=1024-65535 --pass=stdout
```

The first filter identifies anything engaged in port 80 traffic, and the second takes that set and identifies anything that also used an ephemeral port.

A number of IP address helper options are available. `--anyaddress` filters across source and destination addresses simultaneously. `--not-saddress` and `--not-daddress` pass records with addresses that *don't* match the option specification.

Miscellaneous Filtering Options and Some Hacks

`rwfilter` has a couple of direct text output options: `--print-stat` (see [Example 5-11](#)) and `--print-volume-stat`. These can be used to print a summary of the traffic without having to resort to `cut`, `count`, or other display tools. They also will print volumes of records that did *not* pass a filter.

Example 5-11. Using `--print-stat`

```
$ rwfilter --print-volume-stat in/2005/01/07/in-S0_20050107.01 --proto=0-255
```

	Recs	Packets	Bytes	Files
Total	2019	2730488	402105501	1
Pass	2019	2730488	402105501	
Fail	0	0	0	

```
$ rwfilter --print-stat in/2005/01/07/in-S0_20050107.01 --proto=0-255
```

```
Files 1. Read 2019. Pass 2019. Fail 0.
```

Note in [Example 5-11](#) the use of the `--proto=0-255` option. In almost all invocations, `rwfilter` expects *some* form of filtering applied to it, so when you need a filter that passes everything, the easiest approach is just to specify all the protocols. `--print-stat` and `--print-volume-stat` output to `stderr`, so you can still use `stdout` for pass, fail, and all channels.

Like `rwcut`, `rwfilter` has a record limit command. `--max-pass-records` and `--max-fail-records` can be used to limit the number of records passed through a pass or fail channel.

rwfileinfo and Provenance

SiLK filter files contain a fair amount of metadata, which can be accessed using the `rwfileinfo` command (see [Example 5-12](#)). `rwfileinfo` can work with files, as seen in the examples below, or directly on `stdin` by using `stdin` or `-a` as an argument.

Example 5-12. Using `rwfileinfo`

```
$ rwfileinfo in/2005/01/07/in-S0_20050107.01
```

```
in/2005/01/07/in-S0_20050107.01:
```

```
format(id) FT_RWAUGMENTED(0x14)
```

version 2

byte-order littleEndian

compression(id) none(0)

header-length 28

record-length 28

record-version 2

silk-version 0

count-records 2019

file-size 56560

packed-file-info 2005/01/07T01:00:00 ??

\$ rfilter --print-stat in/2005/01/07/in-S0_20050107.01 --proto=6

--pass=example.rwf

Files 1. Read 2019. Pass 1353. Fail 666.

\$ rwfileinfo example.rwf

example.rwf:

format(id) FT_RWGENERIC(0x16)

version 16

byte-order littleEndian

compression(id) none(0)

header-length 156

record-length 52

record-version 5

silk-version 2.1.0

count-records 1353

file-size 70512

command-lines

1 rfilter --print-stat --proto=6 --pass=example.rwf

in/2005/01/07/in-S0_20050107.01

```
$ rfilter --aport=25 example.rwf --pass=example2.rwf --fail=example2_fail.rwf
```

```
$ rwfileinfo example2.rwf
```

```
example2.rwf:
```

```
format(id)    FT_RWGENERIC(0x16)
```

```
version      16
```

```
byte-order    littleEndian
```

```
compression(id) none(0)
```

```
header-length 208
```

```
record-length 52
```

```
record-version 5
```

```
silk-version  2.1.0
```

```
count-records 95
```

```
file-size     5148
```

```
command-lines
```

```
1 rfilter --print-stat --proto=6 --pass=example.rwf
```

```
in/2005/01/07/in-S0_20050107.01
```

```
2 rfilter --aport=25 --pass=example2.rwf
```

```
--fail=example2_fail.rwf example.rwf
```

The fields reported by `rwfileinfo` are as follows:

```
example2.rwf
```

The first line of every `rwfileinfo` dump is the name of the file.

```
format(id)
```

SiLK files are maintained in a number of different optimized formats; the `format` value is a C macro describing the type of the file, followed by the hexadecimal ID of that type.

```
version
```

The version of the file format.

```
byte-order
```

The order in which bytes are stored on disk; SiLK maintains distinct little- and big-endian formats for faster reading.

compression(id)

Whether the file is natively compressed, again for faster reading.

header-length

The size of the file header; a SiLK file with no records will be just the size of the header-length.

record-length

The size of individual file records. This value will be 1 if records are variable length.

record-version

The version of the records (note that record versions are distinct from file versions and SiLK versions).

silk-version

The version of the SiLK suite used to create the file.

count-records

The number of records in the file.

file-size

The total size of the file; if the file is uncompressed, this value should be equivalent to the header length added to the product of the record length and record count.

command-lines

A record of the SiLK commands used to create the file.

Example 5-13 shows how to use the --note-add command.

Example 5-13. Using --note-add

```
$ rfilter --aport=22 example.rwf --note-add='Filtering ssh' --pass=ex2.rwf
```

```
$ rwfileinfo ex2.rwf
```

ex2.rwf:

```
format(id)    FT_RWGENERIC(0x16)
```

```
version      16
```

```
byte-order    littleEndian
```

```
compression(id) none(0)
```

```
header-length 260
```

```
record-length 52
```

```

record-version    5
silk-version      2.1.0
count-records     10
file-size         780
command-lines

    1 rfilter --print-stat --proto=6 --pass=example.rwf
in/2005/01/07/in-S0_20050107.01

    2 rfilter --aport=22 --note-add=Filtering ssh
--pass=ex2.rwf example.rwf
annotations

    1 Filtering ssh

```

Combining Information Flows: **rwcount**

rwcount can produce time series data from the output of an **rfilter** command. It works by placing counts of bytes, packets, and flow records into fixed-duration *bins*, which are equally sized time periods specified by the **rwcount** **bin-size** option. **rwcount** is a relatively straightforward application. Most of its complexity comes from relating the flows, which themselves have a duration, to the bins.

The simplest invocation of **rwcount** is shown in [Example 5-14](#). The first thing to notice is the use of the **bin-size** option. In this example, the bins are half an hour, or 1,800 seconds. If **bin-size** isn't specified, **rwcount** will default to 30-second bins. Bin sizes don't have to be integers; floating-point specifications with a resolution down to the millisecond are acceptable for people who like *lots* of bins in their output.

*Example 5-14. Simple **rwcount** invocation*

```

$ rfilter in/2005/01/07/in-S0_20050107.01 --all=stdout |
  rwcount --bin-size=1800

```

Date	Records	Bytes	Packets
2005/01/07T01:00:00	257.58	42827381.72	248724.14
2005/01/07T01:30:00	1589.61	211453506.60	1438751.93
2005/01/07T02:00:00	171.81	147824612.67	1043011.93

As [Example 5-14](#) shows, **rwcount** outputs four columns: a date column in SiLK's standard date format (YYYY/MM/DDTHH:MM:SS), followed by record, byte, and packet columns. The floating-point values are a function of **rwcount** interpolating how much traffic should be in each bin; **rwcount** calls this a *load scheme*.

The load scheme is an attempt by **rwcount** to approximate how much of a flow took place over the period specified by the bins. In the default load scheme, **rwcount** splits each flow proportionally across all the bins during which the flow was active.

ch the flow was taking place. For example, if a flow takes place from 00:04:00 to 00:11:00, and bins are five minute s long, 1/7 of the flow will be added to the first (00:00:00-00:04:59) bin, 5/7 to the second bin (00:05:00-00:09:59), and 1/7 to the third (00:10:00-00:14:59) bin. `rwcount` takes an integer parameter in the `--load-scheme` option, with t he following results:

0	Split the traffic evenly across all bins covered. In the previous example, the flow would be split into thirds, and a third added to each bin.
1	Add the entire flow to the first bin covered by the flow. In the previous example, 00:00:00-00:04:59.
2	Add the entire flow to the last bin covered by the flow. In the previous example, 00:10:00-00:14:59.
3	Add the entire flow to the middle bin covered by the flow. In the previous example, 00:05:00-00:09:59.
4	The default load scheme.

`rwcount` uses the flow data provided to guess which time bins are required, but sometimes you have to explicitly specify the time, especially when coordinating multiple files. This can be done using the `--start-epoch` and `--end-epoch` options to specify starting and ending bin times. Note that these parameters can use the epoch time or `yyyymm/dd:HH:MM:SS` format. `rwcount` also has an option to print dates using epoch time: the `--epoch-slots` option.

The `--skip-zero` option (see [Example 5-15](#)) is one of a number of output format options. Normally, `rwcount` prints every empty bin it has allocated, but `--skip-zero` causes empty bins to be omitted from the output. In addition, `rwcount` supports many of the output options mentioned for `rwcut`: `--no-titles`, `--no-columns`, `--column-separator`, `--no-final-delimiter`, and `--delimited`.

Example 5-15. Using epoch slots and the `--skip-zero` option

```
rwfilter in/2005/01/07/in-S0_20050107.01 --all=stdout |
  rwcount --bin-size=1800.00 --epoch
    Date|    Records|    Bytes|    Packets|
    1105059600|    257.58|  42827381.72|   248724.14|
    1105061400|   1589.61|  211453506.60|  1438751.93|
    1105063200|    171.81|  147824612.67|   1043011.93|
$ rwfilter in/2005/01/07/in-S0_20050107.01 --all=stdout |
  rwcount --bin-size=1800.00
    --epoch --start-epoch=1105057800
    Date|    Records|    Bytes|    Packets|
```

1105057800	0.00	0.00	0.00
1105059600	257.58	42827381.72	248724.14
1105061400	1589.61	211453506.60	1438751.93
1105063200	171.81	147824612.67	1043011.93

```
$ rfilter in/2005/01/07/in-S0_20050107.01 --all=stdout |
```

```
rwcount --bin-size=1800.00
```

```
--epoch --start-epoch=1105056000
```

Date	Records	Bytes	Packets
1105056000	0.00	0.00	0.00
1105057800	0.00	0.00	0.00
1105059600	257.58	42827381.72	248724.14
1105061400	1589.61	211453506.60	1438751.93
1105063200	171.81	147824612.67	1043011.93

```
$ rfilter in/2005/01/07/in-S0_20050107.01 --all=stdout |
```

```
rwcount --bin-size=1800.00
```

```
--epoch --start-epoch=1105056000 --skip-zero
```

Date	Records	Bytes	Packets
1105059600	257.58	42827381.72	248724.14
1105061400	1589.61	211453506.60	1438751.93
1105063200	171.81	147824612.67	1043011.93

rwset and IP Sets

IP sets are SiLK's most powerful capability, and something that distinguishes the toolkit from most other analytical tools. An IP set is a binary representation of an arbitrary collection of IP addresses. IP sets can be created from text files, from SiLK data, or by using other binary SiLK structures.

The easiest way to start with IP sets is to create one, as in [Example 5-16](#).

Example 5-16. Creating IP sets with rwset

```
$ rfilter in/2005/01/07/in-S0_20050107.01 --all=stdout |
```

```
rwset --sip-file=sip.set --dip-file=dip.set
```

```
$ ls -l *.set
```



```
-rw-r--r-- 1 mcollins staff 580 Jan 10 01:06 dip.set
```

```
-rw-r--r-- 1 mcollins staff 15088 Jan 10 01:06 sip.set
```

```
$ rwsetcat sip.set | head -5
```

```
0.0.0.0
```

```
32.16.40.178
```

```
32.24.41.181
```

```
32.24.215.49
```

```
32.30.13.177
```

```
$ rwfileinfo sip.set
```

```
sip.set:
```

```
format(id)      FT_IPSET(0x1d)
```

```
version         16
```

```
byte-order      littleEndian
```

```
compression(id) none(0)
```

```
header-length   76
```

```
record-length   1
```

```
record-version  2
```

```
silk-version    2.1.0
```

```
count-records   15012
```

```
file-size       15088
```

```
command-lines
```

```
1 rwset --sip-file=sip.set --dip-file=dip.set
```

rwset takes flow records and produces up to four output files. The file specified with `--sip-file` will contain source IP addresses from the flow, `--dip-file` will contain destination addresses, `--any-file` will contain source and destination IP addresses, and `--nhip-file` will contain next hop addresses. The output is binary and read with `rwsetcat`, and as with all SiLK files, the file can be examined using `rwfileinfo`.

The power of IP sets comes when they're combined with `rwfilter`. `rwfilter` has eight commands that accept IP sets (`--sipset`, `--dipset`, `--nhipset`, `--anyset`, and their negations). Sets are explicitly designed so `rwfilter` can rapidly query using them, enabling a variety of useful queries, as seen in [Example 5-17](#).

Example 5-17. Set manipulation and response

\$ # First, we create IP sets; I use a port=123 (NTP on UDP) to filter down

\$ # to a reasonable set of addresses. NTP clients and servers use the same

\$ # port.

\$ rfilter in/2005/01/07/in-S0_20050107.01 --pass=stdout --aort=123 |

rwset --sip-file=sip.set --dip-file=dip.set

\$ # Now, let's see how many IP addresses are created

\$ rwsetcat --count-ip sip.set

15

\$ # Generating output using rfilter; note the use of the --dipset file as the

\$ # sip set; this means that I'm now looking for messages that responded to

\$ # these addresses. This means that I've seen ntp going to and from the

\$ # address, meaning it's likely to be a legitimate speaker, as opposed to a

\$ # scan on port 123.

\$ rfilter out/2005/01/07/out-S0_20050107.01 --dipset=sip.set --pass=stdout

--aort=123 | rwc | head -5

sIP	dIP sPort dPort pro	packets	bytes	\
flags	sTime	dur	eTime sen	
128.3.23.152	56.7.90.229 123 123 17	1	76	\
	2005/01/07T01:10:00.520	0.083	2005/01/07T01:10:00.603	?
128.3.23.152	192.41.221.11 123 123 17	1	76	\
	2005/01/07T01:10:15.519	0.000	2005/01/07T01:10:15.519	?
128.3.23.231	87.221.134.185 123 123 17	1	76	\
	2005/01/07T01:24:46.251	0.005	2005/01/07T01:24:46.256	?
128.3.26.152	58.243.214.183 123 10123 17	1	76	\
	2005/01/07T01:27:08.854	0.000	2005/01/07T01:27:08.854	?

\$ # Let's look at statistics; using the same file, I look at the hosts

\$ # that responded

\$ rfilter out/2005/01/07/out-S0_20050107.01 --dipset=sip.set --aort=123

```
--print-stat
```

```
Files  1. Read  12393. Pass   21. Fail  12372.
```

```
$ # Now I look at everyone else; not-dipset means that I'm looking at everything
```

```
$ # on port 123 that doesn't go to these addresses.
```

```
$ rfilter out/2005/01/07/out-S0_20050107.01 --not-dipset=sip.set --aport=123
```

```
--print-stat
```

```
Files  1. Read  12393. Pass  337. Fail  12056.
```

Sets can also be generated by hand using `rwsetbuild`, which takes text input and produces a set file as the output. The `rwsetbuild` specification takes any of the IP address specifications used by the `--saddress` option in `rfilter`: literal addresses, integers, ranges within dotted quads, and netmasks. [Example 5-18](#) demonstrates this.

Example 5-18. Building a set using `rwsetbuild`

```
$ cat > setsample.txt
```

```
# Comments in set files are prefaced with a hashmark
```

```
# Literal address
```

```
255.230.1.1
```

```
# Note that I'm putting addresses in some semi-random order; the output
```

```
# will be ordered.
```

```
111.2.3-4.1-2
```

```
# Netmask
```

```
22.11.1.128/30
```

```
^D
```

```
$ rwsetbuild setsample.txt setsample.set
```

```
$ rsetcat --print-ip setsample.set
```

```
22.11.1.128
```

```
22.11.1.129
```

```
22.11.1.130
```

```
22.11.1.132
```

```
111.2.3.1
```

```
111.2.3.2
```

111.2.4.1

111.2.4.2

255.230.1.1

Sets can also be manipulated using the `rwsettool` command, which provides a variety of mechanisms for adding and removing sets. `rwsettool` supports four manipulations:

`--union`

Creates a set that includes any address that appears in any of the sets.

`--intersect`

Creates a set that includes only addresses that appear in all the sets specified.

`--difference`

Removes addresses in the latter sets from the first set.

`--sample`

Randomly samples a set to produce a subset.

`rwsettool` is generally invoked using an output path (`--output=_file_`), but if nothing is specified, it will dump to `stdout`. As with `rwfilter`, `rwsettool` output is binary, so a pure terminal dump triggers an error. [Example 5-19](#) shows a manipulation with `rwsettool`.

Example 5-19. Set manipulation with `rwsettool`

```
$ rm setsample2.set
```

```
$ cat > setsample2.txt
```

```
# Build a set that covers our original setsample file to
```

```
# see what happens with various functions
```

```
22.11.1.128/29
```

```
$ rwsetbuild setsample2.txt setsample2.set
```

```
$ rwsettool --union setsample.set setsample2.set | rwsetcat
```

```
22.11.1.128
```

```
22.11.1.129
```

```
22.11.1.130
```

```
22.11.1.131
```

```
22.11.1.132
```

22.11.1.133

22.11.1.134

22.11.1.135

111.2.3.1

111.2.3.2

111.2.4.1

111.2.4.2

255.230.1.1

```
$ rwsettool --intersect setsample.set setsample2.set | rwsetcat
```

22.11.1.128

22.11.1.129

22.11.1.130

22.11.1.131

```
$ rwsettool --difference setsample.set setsample2.set | rwsetcat
```

111.2.3.1

111.2.3.2

111.2.4.1

111.2.4.2

255.230.1.1

rwuniq

rwuniq is the utility knife of counting tools. It allows an analyst to specify a key containing one or more fields, and will then count a number of different values, including total number of bytes, packets, flow records, or unique IP addresses matching the key.

rwuniq's default configuration counts the number of flows that occurred for a particular key. The key itself must be specified using the `--field` option, which accepts the field specifiers in [Table 5-1](#). **rwuniq** can accept multiple fields, and the key will be generated in the order specified in the command line. [Example 5-20](#) demonstrates the key features of the `--field` option. As it shows, field order in the option affects field ordering in the output.

*Example 5-20. Various field specifiers using **rwuniq***

```
$ rwfilter out/2005/01/07/out-S0_20050107.01 --all=stdout |
```

```
rwuniq --field=sip,proto | head -4
```

```

sIP|pro| Records|
131.243.142.85| 17|    1|
131.243.141.187| 17|    6|
128.3.23.41| 17|    4|

$ rwfilter out/2005/01/07/out-S0_20050107.01 --all=stdout |
  rwuniq --field=1,2 | head -4

sIP|    dIP| Records|
128.3.174.158| 128.3.23.44|    2|
128.3.191.1|239.255.255.253|    8|
128.3.161.98|131.243.163.206|    1|

$ rwfilter out/2005/01/07/out-S0_20050107.01 --all=stdout |
  rwuniq --field=sip,sport | head -4

sIP|sPort| Records|
131.243.63.143|53504|    1|
131.243.219.52|61506|    1|
131.243.163.206| 1032|    1|

$ rwfilter out/2005/01/07/out-S0_20050107.01 --all=stdout |
  rwuniq --field=sport,sip | head -4

sPort|    sIP| Records|
55876| 131.243.61.70|    1|
51864|131.243.103.106|    1|
50955| 131.243.103.13|    1|

```

Also, note that when fields' orders are changed, the order in which records are output also changes. `rwuniq` does *not* guarantee record ordering by default; sorting can be ordered by using the `--sort-output` option.

`rwuniq` provides a number of count switches that instruct it to count additional values (see [Example 5-21](#)). The counting switches are `--bytes`, `--packets`, `--flows`, `--sip-distinct`, and `--dip-distinct`. Each of these fields can be used on their own, or by specifying a threshold (e.g., `--bytes`, `--bytes=10`, or `--bytes=10-100`). A single-value threshold (`--bytes=10`) provides a minimum, while a two-value threshold (`--bytes=10-100`) provides a range with a minimum and maximum. If you don't specify an argument, then the switch returns all values.

Example 5-21. Field spec with `rwuniq`

```
$ rfilter out/2005/01/07/out-S0_20050107.01 --all=stdout |
```

```
rwuniq --field=sport,sip --bytes --packets | head -5
```

sPort	sIP	Bytes	Packets
55876	131.243.61.70	308	4
51864	131.243.103.106	308	4
50955	131.243.103.13	308	4
56568	128.3.212.145	360	5

```
$ rfilter out/2005/01/07/out-S0_20050107.01 --all=stdout |
```

```
rwuniq --field=sport,sip --bytes --packets=8 | head -5
```

sPort	sIP	Bytes	Packets
0	131.243.30.224	2520	30
959	128.3.215.60	876	19
2315	131.243.124.237	608	8
56838	131.243.61.187	616	8

```
$ rfilter out/2005/01/07/out-S0_20050107.01 --all=stdout |
```

```
rwuniq --field=sport,sip --bytes --packets=8-20 | head -5
```

sPort	sIP	Bytes	Packets
959	128.3.215.60	876	19
2315	131.243.124.237	608	8
56838	131.243.61.187	616	8
514	128.3.97.166	2233	20

rwbag

The last set of tools to discuss in this chapter are *bag tools*. A *bag* is a form of storage structure. It contains a key (which can be an IP address, a port, the protocol, or an interface index), and a count of values for that key. Bags can be created from scratch or from flow data using the `rwbag` command (see [Example 5-22](#)).

Example 5-22. An `rwbag` call, creating an IP address bag

```
$ rfilter out/2005/01/07/out-S0_20050107.01 --all=stdout |
```

```
rwbag --sip-bytes=sip_bytes.bag
```

```
$ rwbagcat sip_bytes.bag | head -5
```

128.3.2.16	10026403
128.3.2.46	27946
128.3.2.96	218605
128.3.2.98	636
128.3.2.102	1568

Like sets, bags are a second-order binary structure for SiLK, meaning that they have their own toolkit (`rwbagcat`, `rwbagtool`, and `rwbagbuild`), the data is binary (so it can't be read with `cat` or a text editor), and they can be derived from flow data or built from a datafile.

The basic bag generation tool is `rwbag`, which as seen in [Example 5-22](#), takes flow data and produces a bag file from it. `rwbag` can generate 27 types of bags, simultaneously if you're so inclined. These 27 types comprise three types of counting (bytes, packets, and flows), and nine types of key (sip, dip, sport, dport, proto, sensor, input, output, nhop). Combine the key and the counting type, and you have a switch that will create a bag. For example, to count all packets from source and destination IP addresses, call `rwbag --sip-packets=b1.bag --dip-packets=b2.bag`.

Advanced SiLK Facilities

In this section, we discuss more advanced SiLK facilities, in particular, the use of PMAPs and the collection and conversion of SiLK data.

pmaps

A SiLK *prefix map* (PMAP) is a binary file that associates specific subnetworks (prefixes) with tags. PMAPs are used to record various mappings of a network, such as whether a network belongs to a particular organization or ASN, or country code lookup. Using a source such as [GeoIP](#), you can build a PMAP that associates IP addresses with their country of origin.

The SiLK tool suite expects some basic PMAPs:

`address_types.pmap`

Describes an address's type, conventionally indicating whether the address is inside or outside of the network you are monitoring. Specify the default filesystem location for this PMAP using the `SILK_ADDRESS_TYPES` environmental variable.

`country_codes.pmap`

This PMAP describes the country code for an address. Specify the default location for this PMAP using the `SILK_COUNTRY_CODES` environmental variable.

PMAPs, like set files, can be created from text. [Example 5-23](#) shows a simple PMAP file. Note the following attributes:

§ The set of labels at the beginning. PMAPs do not store strings, but enumerable types identified by an integer. This enumeration is defined using the labels. You can see that the PMAP in [Example 5-23](#), for instance, stores a 3 to mark normal traffic.

§ The default key. Any value that doesn't match one of the network blocks listed in the map is given the default value.

§ The actual declarations. Each declaration consists of a network specification, such 192.168.0.0/16, followed by a label.

Example 5-23. PMAP Input

```
# This is a simple PMAP file that tracks some of the standard RFC 1918

# reserved addresses

#

# First we create some labels

label 0 1918-reserved

label 1 multicast

label 2 future

label 3 normal

#

# Specify the mode; this must be either ip or proto-port. ip in this case

# refers to v4 addresses

#

mode ip

#

# Everything otherwise not specified is normal

default normal

# Now the maps

192.168.0.0/16 1918-reserved

10.0.0.0/8 1918-reserved

172.16.0.0/12 1918-reserved

224.0.0.0/4 multicast

240.0.0.0/4 future
```

Once you've created a text representation of the PMAP, you can compile the binary PMAP file using the `rwpm` `apbuild` command. `rwpm` `apbuild` has two mandatory arguments: an input filename, with the file in the text format described above, and a name for the output file. As with most SiLK commands, `rwpm` `apbuild` will not overwrite an existing output file. For example:

```
$ rwpm apbuild -i reserve.txt -o reserve.pmap
```

```
$ ls -l reserve.*
```

```
-rw-r--r-- 1 mcollins staff 406 May 27 17:16 reserve.pmap
```

```
-rw-r--r-- 1 mcollins staff 526 May 27 17:00 reserve.txt
```

Once a PMAP file is created, it can be added to `rwfilter` and `rwcut` using the `pmap-file` argument. Specifying the use of a PMAP file effectively creates a new set of fields in the filter and cut commands; since PMAP files are explicitly related to IP addresses, these new fields are bound to IP addresses.

Consider [Example 5-24](#), which uses `rwcut`. In this example, the `--pmap-file` argument is colon-delimited; the value before the colon (`reserve` in the example) is a label, and the value after is a filename. `rwcut` binds the term `reserve` to the pmaps for the source and destination IP address, creating two new fields: `src-reserve` (for the mapping of the source address to the PMAP) and `dst-reserve` (for the mapping of the destination address) to the PMAP.

Example 5-24. Creating the `src-reserve` and `dst-reserve` fields

```
$ rwcut --pmap-file=reserve:reserve.pmap --fields=1-4,src-reserve,dst-reserve
```

```
traceroute.rwf | head -5
```

```

sIP|      dIP|sPort|dPort| src-reserve| dst-reserve|
192.168.1.12| 192.168.1.1|65428| 53| 1918-reserved| 1918-reserved|
192.168.1.12| 192.168.1.1|56126| 53| 1918-reserved| 1918-reserved|
192.168.1.12| 192.168.1.1|52055| 53| 1918-reserved| 1918-reserved|
192.168.1.1| 192.168.1.12| 53|56126| 1918-reserved| 1918-reserved|
```

\$ # Using the pmap in filter; note that `rwcut` is not using the pmap

```
$ rwfilter --pmap-file=reserve:reserve.pmap --pass=stdout traceroute.rwf
```

```
--pmap-src-reserve=1918-reserved | rwcut --field=1-5
```

```
| head -5
```

```

sIP| dIP|sPort|dPort|pro|
192.168.1.12| 192.168.1.1|65428| 53| 17|
192.168.1.12| 192.168.1.1|56126| 53| 17|
192.168.1.12| 192.168.1.1|52055| 53| 17|
192.168.1.1| 192.168.1.12| 53|56126| 17|
```

Collecting SiLK Data

There are a number of different tools for collecting data and pushing it into SiLK. The major ones are YAF, which is a flow collector, and `rwptoflow` and `rwcut`, which convert other data into SiLK format.

YAF

Yet Another Flowmeter (YAF) is the reference implementation for the IETF IPFIX standard, and is the standard flow collection software for the SiLK toolkit. YAF can read *pcap* data from files or capture packets directly, which it then assembles into flow records and exports to disk. It has [online documentation](#). The tool itself can be entirely configured using command-line options, but the number of options is fairly daunting. At its simplest, a YAF command looks like this:

```
$ sudo yaf -i en1 --live=pcap -out /tmp/yaf/yaf
```

This reads data from interface *en1* and drops it to the file in the temporary directory. Additional options control how data is read and how it is converted into flow and output format.

YAF output is specified via the `--out` switch in tandem with the `--ipfix` and `--rotate` switches. By default, `--out` outputs to a file; in the example above, the file is `/tmp/yaf/yaf`, but any valid filename will do (if `--out` is set to `-`, then YAF will output to `stdout`).

When `--out` is specified with `--rotate`, YAF writes the output to files that are rotated by a delay specified by the `--rotate` switch (e.g., `--rotate 3600` will update files every hour). In this mode, YAF uses the name specified by `--out` as a base filename, and attaches a suffix specified in `YYYYMMDDhhmmss` format, along with a decimal serial number and then a `.yaf` file extension.

When YAF is specified with the `--ipfix` switch, it communicates IPFIX data to a daemon located elsewhere on the network. In this case (the most complicated option), `--ipfix` takes a transport protocol as an argument, while `--out` takes the IP address of the host. The additional `--ipfix-port` switch takes a port number when needed. Consult the documentation for more information.

The most important options are:

`--live`

Specifies the type of data being read; possible values are `pcap`, `dag`, or `napatech`. `dag` and `napatech` refer to proprietary packet capture systems, so unless you have that hardware, just set `--live` to `pcap`.

`--filter`

Applies a BPF filter to the *pcap* data.

`--out`

The output specifier, discussed above. The output specifier will be a file, a file prefix, or an IP address depending on whatever other switches are used.

`--ipfix`

Takes a transport protocol (`tcp`, `udp`, `sctp`, or `spread`) as an argument, and specifies that output is IPFIX transported over the network. Consult the YAF documentation for more information.

`--ipfix-port`

Used only if `--ipfix` is specified. It specifies the port that the IPFIX data is sent to.

`--rotate`

Used only with files. If present, the filename in `--out` is used as a prefix, and files are written with a timestamp appended to them. The `--rotate` option takes an argument and the number of seconds before moving to a new file.

--silk

Specifies output that can be parsed by SiLK's rflowpack tools.

--idle-timeout

Specifies the idle timeout for flows in seconds. If a flow is present in the flow cache and isn't active, it's flushed as soon as it's been inactive for the duration of the idle timeout. Defaults to 300 seconds (five minutes).

--active-timeout

Specifies the active timeout for flows; the active timeout is the maximum amount of time an active flow will be stored in cache before being flushed. Defaults to 30 minutes (1,800 seconds). Note that the active timeout determines the maximum observed duration of collected flows.

YAF has many more options, but these are the basic ones to consider when configuring flows. Consult the YAF manpage for more details.

COOKBOOK: YAF

YAF has a ton of options, and how they operate together can be a bit confusing. Here are some examples of YAF invocations:

Read yaf from an interface (en1) and write to a file on disk:

```
$sudo yaf -i en1 --live=pcap -o /tmp/yaf/yaf
```

Rotate the files every five minutes:

```
$sudo yaf -i en1 --rotate 300 --live=pcap -o /tmp/yaf/yaf
```

Read a file from disk and convert it:

```
$yaf <example.pcap >yafout
```

Run a BPF filter on the data, in this case for TCP data only

```
$ sudo yaf -i en1 --rotate 300 --live=pcap -o /tmp/yaf/yaf --filter="tcp"
```

Export the YAF data over IPFIX to address 128.2.14.11:3059

```
$ sudo yaf --live pcap --in eth1 --out 128.2.14.11 --ipfix-port=3059
```

```
--ipfix tcp
```

rwptoflow

SiLK uses its own compact binary formats to represent NetFlow data that tools such as rwcut and rwcount present in a human-readable form. There are times when an analyst needs to convert other data into SiLK format, such as taking packet captures from IDS alerts and converting it into a format where IP set filtering can be done on the data.

The go-to tool for this task is rwptoflow. rwptoflow is a packet data to flow conversion tool. It does *not* aggregate flows; instead, each flow generated by rwptoflow is converted into a one-packet flow record. The resulting file c

an then be manipulated by the SiLK suite as any other flow file.

rwptoflow is invoked relatively simply with an input filename as its argument. In [Example 5-25](#), the *pcap* data from a traceroute is converted into flow data using *rwptoflow*. The resulting raw file is then read using *rwcut* and you can see the correspondence between the traceroute records and the resulting flow records.

Example 5-25. Converting pcap data with rwptoflow

```
$ tcpdump -v -n -r traceroute.pcap | head -6

reading from file traceroute.pcap, link-type EN10MB (Ethernet)

21:06:50.559146 IP (tos 0x0, ttl 255, id 8010, offset 0, flags [none],
    proto UDP (17), length 64)

    192.168.1.12.65428 > 192.168.1.1.53: 63077+ A? jaws.oscar.aol.com. (36)

21:06:50.559157 IP (tos 0x0, ttl 255, id 37467, offset 0, flags [none],
    proto UDP (17), length 86)

    192.168.1.12.56126 > 192.168.1.1.53: 30980+ PTR?
    dr._dns-sd._udp.0.1.168.192.in-addr.arpa. (58)

21:06:50.559158 IP (tos 0x0, ttl 255, id 2942, offset 0, flags [none],
    proto UDP (17), length 66)

    192.168.1.12.52055 > 192.168.1.1.53: 990+ PTR? db._dns-sd._udp.home. (38)

$ rwptoflow traceroute.pcap > traceroute.rwf

$ rwcut --num-recs=3 --fields=1-5 traceroute.rwf

sIP| dIP|sPort|dPort|pro|
192.168.1.12| 192.168.1.1|65428| 53| 17|
192.168.1.12| 192.168.1.1|56126| 53| 17|
192.168.1.12| 192.168.1.1|52055| 53| 17|
```

rwcut

When correlating data between different sources, you will occasionally want to convert it into SiLK's format. *rwcut* is the default tool for converting data into SiLK representation, as it works with columnar text files. Using *rwcut*, you can convert IDS alerts and other data into SiLK data for further manipulations.

The easiest way to invoke *rwcut* is to use it as an inverse of *rwcut*. Create a file with columnar entries and make sure that the titles match those used by *rwcut*:

```
$cat rwcut_sample.txt
```

```

sIP    |dIP    |proto
128.2.11.4 | 29.3.11.4 | 6
11.8.3.15  | 9.12.1.4  | 17
$ rwtuc < rwtuc_sample.txt > rwtuc_sample.rwf
$ rwcut rwtuc_sample.rwf --field=1-6

sIP| dIP|sPort|dPort|pro|  packets|
128.2.11.4| 29.3.11.4| 0| 0| 6|    1|
11.8.3.15| 9.12.1.4| 0| 0| 17|   1|

```

As the following fragment shows, `rwtuc` will read the columns, use the headers to determine column content, and stuff any unspecified fields with a default value if no column is provided. `rwtuc` can also take column specifications at the command line using the `--fields` and `--column-separator` switches, as so:

```

$cat rwtuc_sample2.txt

128.2.11.4 x 29.3.11.4 x 6 x 5
7.3.1.1   x 128.2.11.4 x 17 x 3

$ rwtuc --fields=sip,dip,proto,packets --column-sep=x < rwtuc_sample2.txt
> rwtuc_sample2.rwf

$ rwcut --fields=1-7 rwtuc_sample2.rwf

sIP| dIP|sPort|dPort|pro|  packets|  bytes|
128.2.11.4| 29.3.11.4| 0| 0| 6|    5|    5|
7.3.1.1| 128.2.11.4| 0| 0| 17|   3|    3|

```

SiLK's binary format requires values for every field, which means that `rwtuc` makes a best guess for field values that it doesn't have. For instance, the previous example specifies packets as a field but not bytes, so `rwtuc` just defines the packet value to be identical to the byte value.

If there exists a common default value (e.g., all traffic has the same protocol), this value can be defined using one of a number of field-stuffing options in `rwtuc`. These options are identical to the field filtering options in `rwfilt`, except they only take single values. For example, `--proto=17` sets the protocol of every entry to 17.

In the fragment below, we use the field stuffing command `--bytes=300` to set a value of 300 bytes for every entry in `rwtuc_sample2.txt`:

```

$ rwtuc --fields=sip,dip,proto,packets --column-sep=x --bytes=300 <
rwtuc_sample2.txt > rwtuc_sample2.rwf

$ rwcut --fields=1-7 rwtuc_sample2.rwf

```

```
sIP| dIP|sPort|dPort|pro|  packets|  bytes|
```

```
128.2.11.4| 29.3.11.4| 0| 0| 6|    5|    300|
```

```
7.3.1.1| 128.2.11.4| 0| 0| 17|   3|    300|
```

The resulting RWF file will contain a value of 300 bytes, even though the byte value is not in the original text file. The packet values, which are specified in the file, are set to whatever was specified there.

Further Reading

1. Time Shimeall, Sid Faber, Markus DeShon, and Drew Kompanek, "Using SiLK for Network Traffic Analysis," Software Engineering Institute.

^[5] You'll notice that there are two datasets, one with scans and one without. To understand why, read Pang *et al.*, "The Devil and Packet Trace Anonymization," ACM CCR 36(1), January 2006.

